

# CS420: Operating Systems

## Threads

---

James Moscola

Department of Engineering & Computer Science

York College of Pennsylvania



# Threads

---

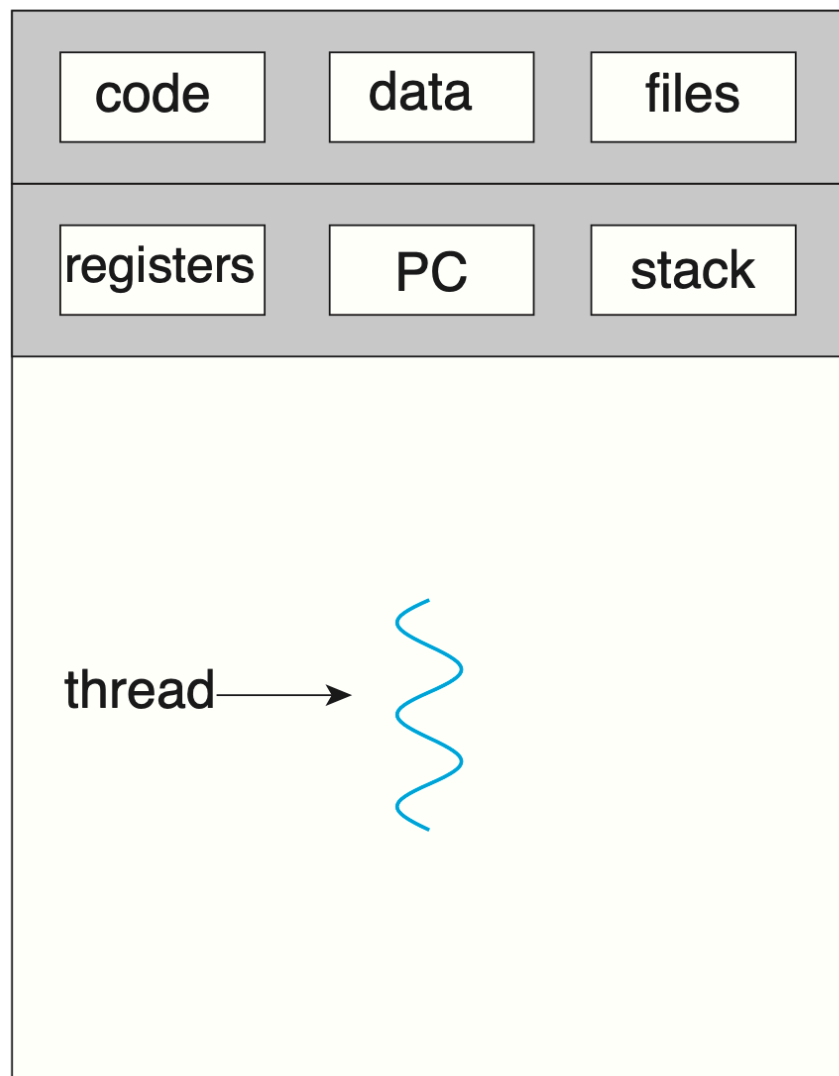
- **A thread is a basic unit of processing**
  - Has the following components:
    - Thread ID
    - Program counter
    - Register set
    - Stack
  - Shares some resources with other threads in same process
    - Code section
    - Data section
    - OS Resources (e.g. open files, signals, etc.)

# Threads

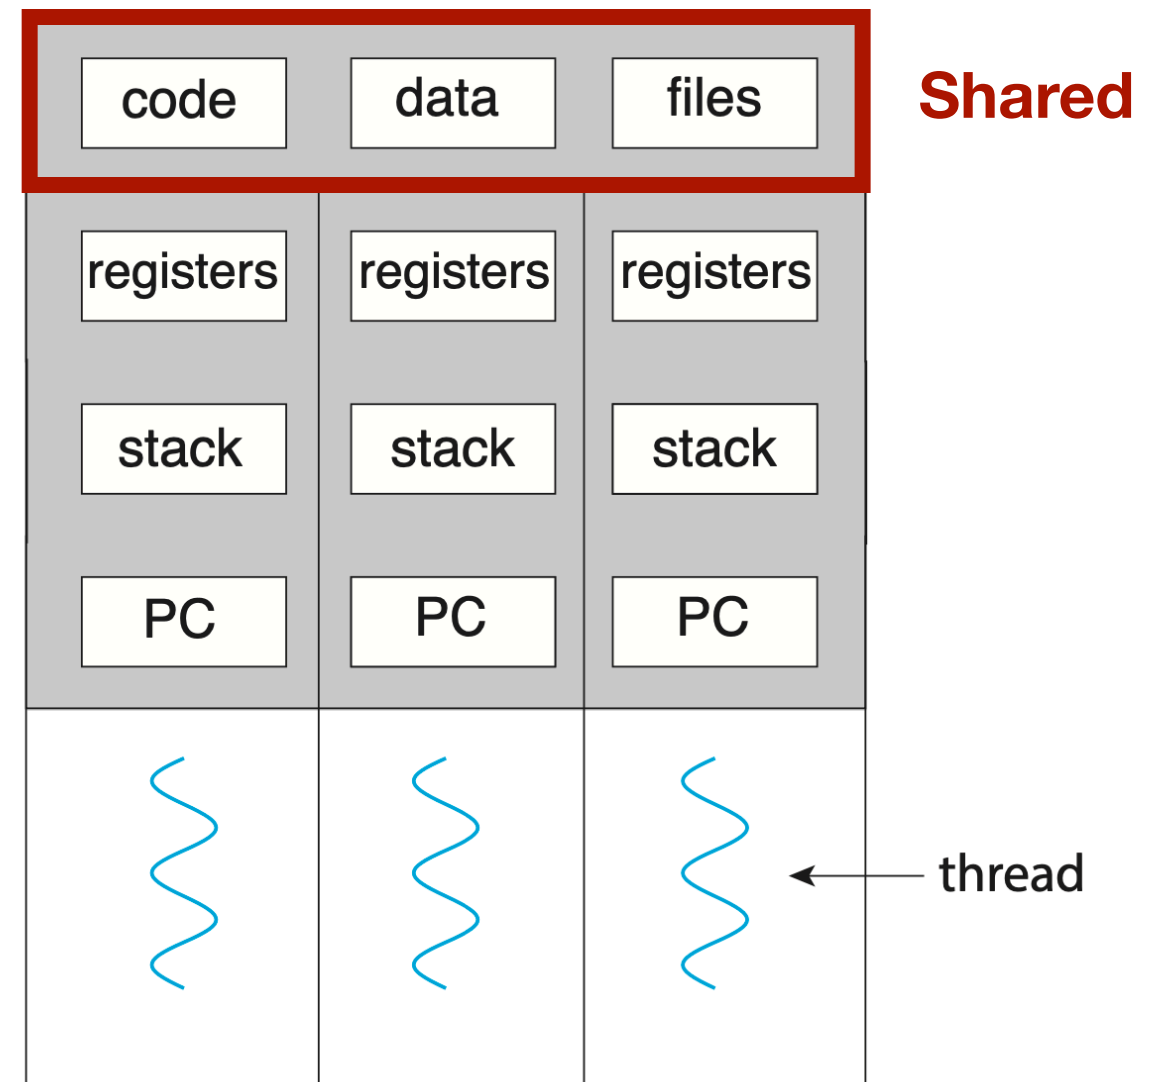
---

- **A heavyweight process is a process that has a single thread of control**
  - Can only perform a single task at a time
- **A multi-threaded process is a process that has multiple threads of control**
  - Can perform more than one task at a time
    - Render images (e.g. thumbnails for an image library)
    - Fetch data
    - Update display
    - Check spelling

# Single and Multithreaded Processes



single-threaded process



multithreaded process

# Thread vs Process

---

- **Processes**

- Independent units of execution
- Each contains its own state information
- Each contains its own address space
- Interact with each other through various IPC mechanisms

- **Threads (within the same process)**

- Share the same state
- Share the same memory space
- Share the same variables
- Can communicate directly through shared variables
- Share signal handling

# Benefits of Multithreaded Programming

---

- **Responsiveness**

- Interactive applications are more responsive when multithreaded (e.g. a thread for the GUI, another for socket, a third for rendering, etc.)

- **Resource Sharing**

- Unlike processes, threads share memory (code and data sections) and resources

- **Economy**

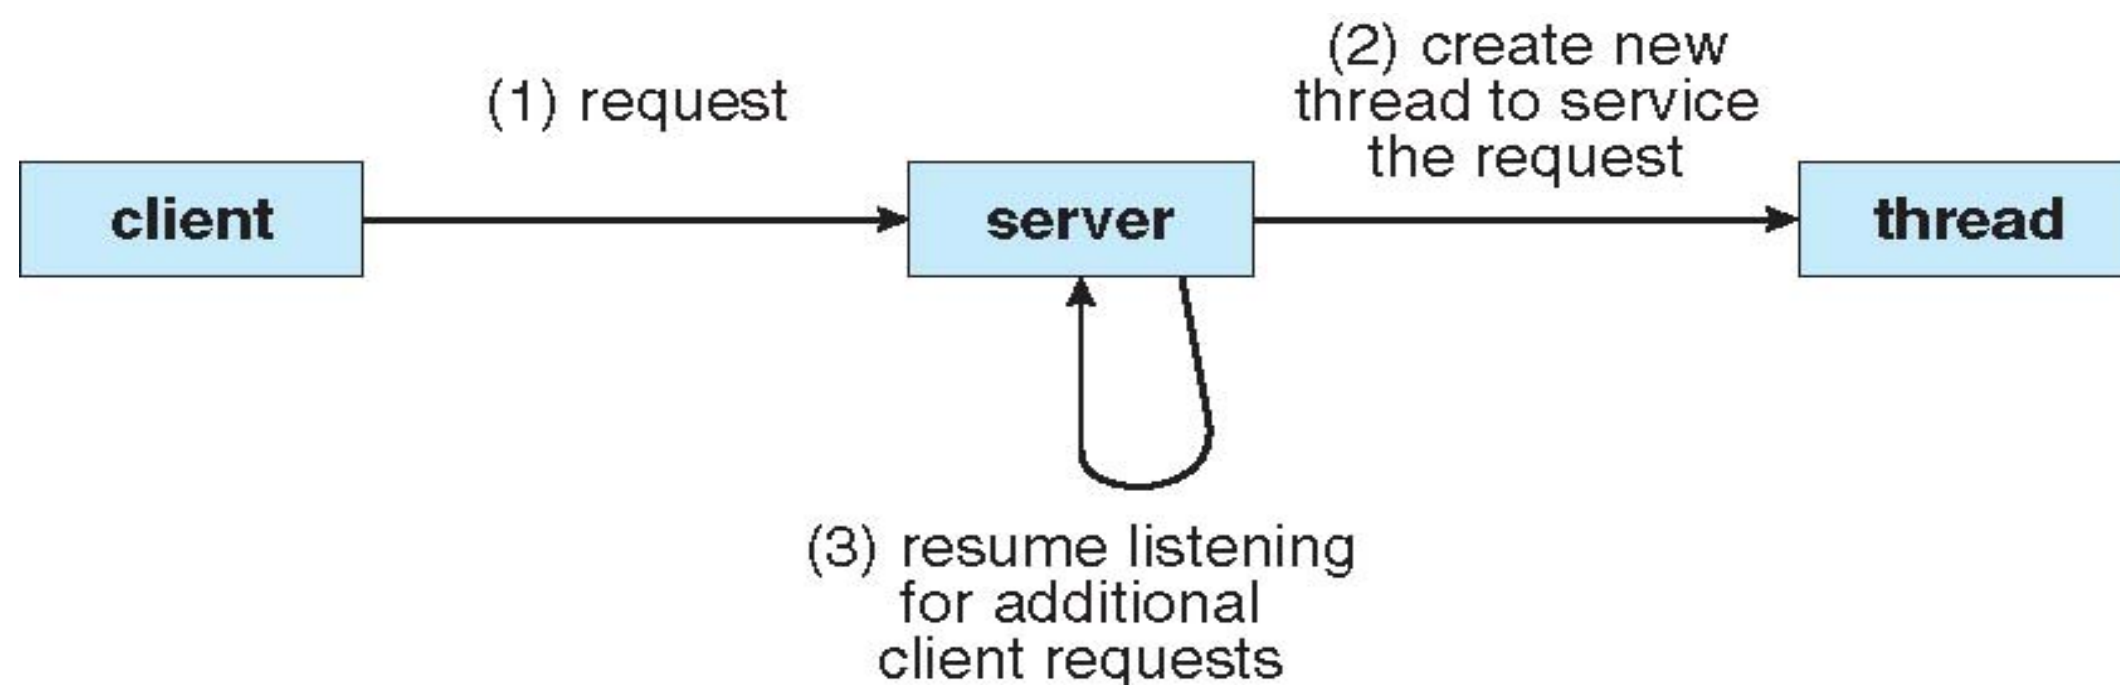
- Since threads share resources, creating threads and switching between them is more efficient than processes

- **Scalability**

- Multithreading allows for increased parallelism on multicore systems as each thread can run on a different CPU core (kernel threads only)


# Multithreaded Server Architecture

---



# Concurrent Execution on a Single-core System

---




- Only a single thread can execute at a time
- Threads are interleaved so each gets time on the processor



# Parallel Execution on a Multicore System

---



- With multiple cores, threads can be divided over the cores and run in parallel
- May still interleave threads if not enough cores are available for all of the threads


# Multicore Programming

---


- **The performance of a process can be increased by properly threading the process to take advantage of modern multicore CPUs**
- **Multicore systems have challenges not faced in a single-core/single-threaded environment**
  - **Dividing activities** - How can an application be divided into separate, concurrent tasks?
  - **Balance** - How can those tasks be divided in such a way that each does an equal amount of work?
  - **Data splitting** - Can the data for those task be divided for processing on separate CPU cores?
  - **Data dependency** - Are there data dependencies between different tasks?
  - **Testing and debugging** - What is the best way to debug a multithreaded program with many different execution paths?

# Data Parallelism and Task Parallelism

- **Data Parallelism** - distribute subsets of the same data across multiple cores



- **Task Parallelism** - distribute threads across multiple cores (each performing a unique operation)



# Thread Support

---

- **Threads may be supported at different levels of the OS**
  - **User threads**
    - Supported above the kernel
    - Managed and scheduled without kernel support
    - Main user thread libraries currently in use  
POSIX PThreads / Java threads
  - **Kernel threads**
    - Supported by the kernel/operating system
    - Managed and scheduled by the kernel/operating system
    - Most modern operating systems support kernel threads  
(e.g. Windows 2000/XP/.../10/11, Solaris, Linux, macOS)


# Multithreading Models

---

- **There must be a relationship between user level threads and kernel threads**
- **Different models of threading exist to define this relationship**
  - Many-to-One
  - One-to-One
  - Many-to-Many

# Many-to-One

- **Many user-level threads mapped to single kernel thread**
  - Examples: Solaris Green Threads, GNU Portable Threads
  - Not many systems use this model
- **Thread management is done in user space**
- **Entire process will block if any single thread blocks (no other threads will run)**
- **Unable to run multiple user-level threads in parallel on a multiprocessor system**
  - Not very common anymore



# One-to-One

- **Each user-level thread maps to kernel thread**


- Examples: Windows NT/XP/2000, Linux, Solaris 9 and later

- **Allows more concurrency**

- A thread can run when another thread has made a blocking system call
- Multiple user-level threads can run in parallel on multiprocessor systems


- **Downside — for each thread created, a corresponding kernel thread must also be created**

- Large number of kernel threads may degrade system performance



# Many-to-Many Model


- **Many user-level threads are mapped to a smaller number of kernel threads**
  - Avoids blocking of threads when a single thread makes a blocking system call
- **Allows the operating system to create a sufficient number of kernel threads**
  - OS may allocate more kernel threads on a machine with more CPU cores
- **Reduces the overhead associated with too many kernel threads as was present in the one-to-one model**





# Two-level Model

- **Similar to the many-to-many model except that it allows a user thread to be bound to a specific kernel thread**
- **Examples include**
  - IRIX
  - HP-UX
  - Tru64 UNIX
  - Solaris 8 and earlier



# Thread Libraries

---

- **A thread library provides programmer with API for creating and managing threads**
- **Two primary ways of implementing**
  - Library entirely in user space (no kernel support)
  - Kernel-level library supported by the OS
- **Three main thread libraries currently in use**
  - (1) **POSIX Pthreads** - user-level or kernel-level threads for POSIX-compliant systems
  - (2) **Windows thread library** - kernel-level threads for Windows systems
  - (3) **Java threads** - threads created and managed in Java programs (typically mapped to thread library of host system)

# POSIX Pthreads

---

- **May be provided either as user-level or kernel-level**
- **A POSIX standard (IEEE 1003.1c) API for thread creation and synchronization**
- **API specifies behavior of the thread library, implementation is up to development of the library**
- **Common in UNIX operating systems (Solaris, Linux, macOS)**

# Pthreads Example

---

```
#include <pthread.h>
#include <stdio.h>

#include <stdlib.h>

int sum; /* this data is shared by the thread(s) */
void *runner(void *param); /* threads call this function */

int main(int argc, char *argv[])
{
    pthread_t tid; /* the thread identifier */
    pthread_attr_t attr; /* set of thread attributes */

    /* set the default attributes of the thread */
    pthread_attr_init(&attr);
    /* create the thread */
    pthread_create(&tid, &attr, runner, argv[1]);
    /* wait for the thread to exit */
    pthread_join(tid, NULL);

    printf("sum = %d\n", sum);
}
```

## Pthreads Example (Cont.)

---

```
/* The thread will execute in this function */
void *runner(void *param)
{
    int i, upper = atoi(param);
    sum = 0;

    for (i = 1; i <= upper; i++)
        sum += i;

    pthread_exit(0);
}
```

# Pthreads Code for Joining 10 Threads

---

```
#define NUM_THREADS 10

/* an array of threads to be joined upon */
pthread_t workers[NUM_THREADS];

for (int i = 0; i < NUM_THREADS; i++)
    pthread_join(workers[i], NULL);
```

# Windows API - Multithreaded C Program

---

```
#include <windows.h>
#include <stdio.h>
DWORD Sum; /* data is shared by the thread(s) */
int main(int argc, char *argv[])
{
    DWORD ThreadId;
    HANDLE ThreadHandle;
    int Param;

    Param = atoi(argv[1]);
    /* create the thread */
    ThreadHandle = CreateThread(
        NULL, /* default security attributes */
        0, /* default stack size */
        Summation, /* thread function */
        &Param, /* parameter to thread function */
        0, /* default creation flags */
        &ThreadId); /* returns the thread identifier */

    /* now wait for the thread to finish */
    WaitForSingleObject(ThreadHandle, INFINITE);

    /* close the thread handle */
    CloseHandle(ThreadHandle);

    printf("sum = %d\n", Sum);
}
```

# Windows API - Multithreaded C Program (Cont.)

---

```
/* The thread will execute in this function */
DWORD WINAPI Summation(LPVOID Param)
{
    DWORD Upper = *(DWORD*)Param;
    for (DWORD i = 1; i <= Upper; i++)
        Sum += i;
    return 0;
}
```



# Java Threads

---

- **Java threads are managed by the JVM (Java Virtual Machine)**
- **Typically implemented using the threads model provided by underlying OS**
- **Java threads may be created in two different ways:**
  - Extend the Thread class and override `run()` method
  - Implement the `Runnable` interface

# Java Multithreaded Program

---

```
import java.util.concurrent.*;

class Summation implements Callable<Integer>
{
    private int upper;
    public Summation(int upper) {
        this.upper = upper;
    }

    /* The thread will execute in this method */
    public Integer call() {
        int sum = 0;
        for (int i = 1; i <= upper; i++)
            sum += i;

        return new Integer(sum);
    }
}

public class Driver
{
    public static void main(String[] args) {
        int upper = Integer.parseInt(args[0]);

        ExecutorService pool = Executors.newSingleThreadExecutor();
        Future<Integer> result = pool.submit(new Summation(upper));

        try {
            System.out.println("sum = " + result.get());
        } catch (InterruptedException | ExecutionException ie) { }
    }
}
```

# Implicit Threading

---

- **Creation and management of threads done by compilers and run-time libraries rather than programmers**
  - Growing in popularity as numbers of threads increase, program correctness more difficult with explicit threads
- **Three methods explored**
  - Thread Pools
  - OpenMP
  - Grand Central Dispatch
- **Other methods include Intel Threading Building Blocks (TBB), `java.util.concurrent` package**

# Thread Pools

---

- **In applications where threads are repeatedly being created/destroyed thread pools might provide a performance benefit**
  - Example: A server that spawns a new thread each time a client connects to the system and discards that thread when the client disconnects
- **A thread pool is a group of threads that have been pre-created and are available to do work as needed**
  - Threads may be created when the process starts
  - A thread may be kept in a queue until it is needed
  - After a thread finishes, it is placed back into a queue until it is needed again
  - Avoids the extra time needed to spawn new threads when they're needed

# Thread Pools

---

- **Advantages of thread pools:**

- Typically faster to service a request with an existing thread than create a new thread (performance benefit)
- Bounds the number of threads in a process
  - The only threads available are those in the thread pool
  - If the thread pool is empty, then the process must wait for a thread to re-enter the pool before it can assign work to a thread
  - Without a bound on the number of threads in a process, it is possible for a process to create so many threads that all of the system resources are exhausted

# Implicit Threading with OpenMP

---

- **Set of compiler directives and an API that provides support for parallel programming in shared-memory environments**
- **Identifies parallel regions as blocks of code that may run in parallel**
  - Developers insert compiler directives into their code to define parallel region

# Implicit Threading with OpenMP (Cont.)

---

```
#include <omp.h>
#include <stdio.h>

int main(int argc, char *argv[])
{
    /* standard sequential code here */

    /* the next bit is automatically parallelized */
    #pragma omp parallel
    {
        printf("I am a parallel region.");
    }

    #pragma omp parallel for
    for (int i = 0; i < 1000; i++) {
        c[i] = a[i] + b[i];
    }

    /* more standard sequential code can go here */

    return 0;
}
```

# Grand Central Dispatch

---

- **Developed for macOS / iOS (and other Apple operating systems)**
- **Tasks are placed into dispatch queues as program runs**
- **Utilizes multiple dispatch queues**
  - **Serial** dispatch queue
    - Tasks are removed in FIFO order
    - Once a task is removed, it must be completed prior to next dequeue
  - **Concurrent** dispatch queue
    - Tasks are removed in FIFO order
    - Multiple tasks may be removed at a time and run in parallel