

CS420: Operating Systems

Threading Issues

James Moscola

Department of Engineering & Computer Science

York College of Pennsylvania



Threading Issues

- **There are a variety of issues to consider with multithreaded programming**
 - Semantics of `fork()` and `exec()` system calls
 - Signal handling
 - Synchronous and asynchronous
 - Thread cancellation
 - Asynchronous or deferred
 - Thread-specific data
 - Create facility needed for data private to thread

Semantics of `fork()` and `exec()`

- Recall that when `fork()` is called, a separate, duplicate process is created
- How should `fork()` behave in a multithreaded program?
 - Should all threads be duplicated?
 - Should only the thread that made the call to `fork()` be duplicated?
- In some systems, different versions of `fork()` exist depending on the desired behavior
 - Some UNIX systems have `fork1()` and `forkall()`
 - `fork1()` only duplicates the calling thread
 - `forkall()` duplicates all of the threads in a process
 - In a POSIX-compliant system, `fork()` behaves the same as `fork1()`

Semantics of `fork()` and `exec()`

- **The `exec()` system call continues to behave as expected**
 - Replaces the entire process that called it, including all threads
- **If planning to call `exec()` after `fork()`, then there is no need to duplicate all of the threads in the calling process**
 - All threads in the child process will be terminated when `exec()` is called
 - Use `fork1()`, rather than `forkall()` if using in conjunction with `exec()`

Signal Handling

- **Signals** are used in **UNIX** systems to notify a process that a particular event has occurred
 - CTRL-C is an example of an **asynchronous signal** that might be sent to a process
 - An asynchronous signal is one that is generated from outside the process that receives it
 - Divide by 0 is an example of a **synchronous signal** that might be sent to a process
 - A synchronous signal is delivered to the same process that caused the signal to occur
- **All signals follow the same basic pattern:**
 - A signal is generated by particular event
 - The signal is delivered to a process
 - The signal is handled by a **signal handler** (all signals are handled exactly once)

Signal Handling


- **Signal handling is straightforward in a single-threaded process**
 - The one (and only) thread in the process receives and handles the signal
- **In a multithreaded program, where should signals be delivered?**
 - Options:
 - (1) Deliver the signal to the thread to which the signal applies
 - (2) Deliver the signal to every thread in the process
 - (3) Deliver the signal only to certain threads in the process
 - (4) Assign a specific thread to receive all signals for the process

Signal Handling

- **Option 1 - Deliver the signal to the thread to which the signal applies**
 - Most likely option when handling synchronous signals (e.g. only the thread that attempts to divide by zero needs to know of the error)
- **Option 2 - Deliver the signal to every thread in the process**
 - Likely to be used in the event that the process is being terminated (e.g. a CTRL-C is sent to terminate the process, all threads need to receive this signal and terminate)

Thread Cancellation

- **Thread cancellation** is the act of terminating a thread before it has completed
 - Example - clicking the stop button on your web browser will stop the thread that is rendering the web page
- The thread to be cancelled is called the **target thread**
- Threads can be cancelled in a couple of ways
 - **Asynchronous cancellation** terminates the target thread immediately
 - Thread may be in the middle of writing data ... not so good
 - **Deferred cancellation** allows the target thread to periodically check if it should be cancelled
 - Allows thread to terminate itself in an orderly fashion



Thread-Specific Data

- **Thread-specific data** - in some applications it may be useful for each thread to have its own copy of data
 - May also be referred to as **Thread-local storage** or **Thread-static variables**
 - The `errno` variable is thread-specific

```
// In C#  
class FooBar {  
    [ThreadStatic] static int foo;  
}
```

```
// In Java  
private static ThreadLocal<Integer> threadLocalInt =  
    new ThreadLocal<Integer>();
```

```
// In a POSIX-compliant system  
// see pthread_key_create()  
//      pthread_setspecific()  
//      pthread_getspecific()
```

Thread Examples - Windows Threads

- **Uses the one-to-one thread model to implement threads**
- **Each thread contains**
 - A thread id
 - Register set (including a program counter)
 - Separate user and kernel stacks (for user-mode and kernel-mode activity)
 - Private data storage area
- **Also implements a fiber — a unit of execution that must be manually scheduled by the application**
 - Every thread can have multiple fibers
 - Fibers run in the context of the thread that created them

Linux Threads

- **Linux oftentimes uses the term `task` rather than `process` or `thread`**
 - Doesn't really distinguish between processes and threads
- **Thread creation is done through the `clone()` system call**
 - The `clone()` can create either 'threads' or 'processes' depending on the options passed to `clone()`
 - The options passed to clone determine how much sharing is taking place between the parent and the child
 - A 'process' can still be created using the `fork()` system call
 - Provide a specific set of options and the `clone()` and `fork()` systems calls behave identically

Linux Threads (Cont.)

- The following table shows the various flags that can be passed to `clone` to determine how much sharing is taking place between the parent and the child

flag	meaning
<code>CLONE_FS</code>	File-system information is shared.
<code>CLONE_VM</code>	The same memory space is shared.
<code>CLONE_SIGHAND</code>	Signal handlers are shared.
<code>CLONE_FILES</code>	The set of open files is shared.

- If ALL these flags are passed to `clone()`, the task is equivalent to a thread
 - Shared memory space, shared list of open files, etc.
- If NONE of these flags are passed to `clone()`, it behaves more like `fork()`
 - NO shared memory space, NO shared list of open files, etc.